

2 Sentiment Analysis

2.2 Movie Review Data

Let us first start by looking at the data provided with the exercise. We have positive and negative movie reviews labeled by human readers, all positive and negative reviews are in the 'pos' and 'neg' folders respectively. If you look inside a sample file, you will see that these review messages have been 'tokenized', where all words are separated from punctuations. There are approximately 1000 files in each category with file names starting with cv000, cv001, cv002 and so on. You will split the dataset into training set and testing set.

1. Write some code to load the data from text files.

Load data from text files

1. % glob เพื่ออ่านชื่อไฟล์ทั้งหมดใน directory
2. % with open เพื่อเปิดไฟล์ text files ที่อยู่ใน directory แล้ว append ไปยัง raw_data

```
In [4]: #Import Libraries
import pandas as pd
import numpy as np
import glob

In [5]: neg_txt = glob.glob( r'C:\Users\dell\Desktop\module8&9\AI\Lab02\neg\*.txt') #read file name in folder' neg
pos_txt = glob.glob( r'C:\Users\dell\Desktop\module8&9\AI\Lab02\pos\*.txt') #read file name in folder' pos

In [6]: # read data in txt.file and append it to list
raw_data = []
for i in range(len(pos_txt)):
    with open(pos_txt[i], 'r') as F:
        raw_data.append(F.read())

for i in range(len(neg_txt)):
    with open(neg_txt[i], 'r') as F:
        raw_data.append(F.read())
```

2.3 TF-IDF

From a raw text review, you want to create a vector, whose elements indicate the number of each word in each document. The frequency of all words within the documents are the 'features' of this machine learning problem.

A popular method for transforming a text to a vector is called tf-idf, short for term frequencyinverse document frequency.

1. Conduct a research about tf-idf and explain how it works.
2. Scikit-learn provides a module for calculating this, this is called TfidfVec- torizer. You can study how this function is used here:

http://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.TfidfVectorizer.html

Write code to transform your text to tf-idf vector.

TF-idf

จะนับความถี่ frequency ของ word ที่อยู่ใน raw_data ทั้งหมด แล้ว return ออกมาเป็น column ซึ่งขนาดของ column จะเท่ากับ จำนวน word ที่ไม่ซ้ำกันทั้งหมด โดยในแต่ละ column ค่าออกมาอยู่ในช่วง 0-1

1. โดยถ้าค่าใกล้ 1 หมายถึง มี word นั้นอยู่ใน text.file นั้นบ่อยมาก
2. ถ้าค่าใกล้ 0 หมายถึง มี word นั้นอยู่ใน text.file น้อยมาก

```
In [7]: from sklearn.feature_extraction.text import TfidfVectorizer
Tfidf = TfidfVectorizer()
feature1 = Tfidf.fit_transform(raw_data)

In [8]: data = pd.DataFrame(feature1.toarray())

In [17]: data.head()
```

```
In [17]: data.head()
```

```
Out[17]:
```

	0	1	2	3	4	5	6	7	8	9	...	39649	39650	39651	39652	39653	39654	39655	39656	39657	39658
0	0.052257	0.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.000000	0.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2	0.000000	0.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	0.000000	0.020881	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
4	0.000000	0.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

5 rows × 39659 columns

```
In [5]: # assign true class
y = np.hstack((np.ones(1000),np.zeros(1000))) #1 = pos , 0 = neg
```

2.4 Classification

Use 4 different models to classify each movie into positive or negative category.

1. K-Nearestneighbor model, using module `sklearn.neighbors.KNeighborsClassifier`
2. RandomForest, using module `sklearn.ensemble.RandomForestClassifier`
3. SVM, using module `sklearn.svm.SVC`
4. Neural network, using `sklearn.neural_network.MLPClassifier`

You may pick other models you would like to try. Just present results for at least 4 models. Please provide your code for model fitting and cross validation. Calculate your classification accuracy, precision, and recall.

```
In [6]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
```

```
In [7]: KN1 = KNeighborsClassifier()
yhat_KN1 = cross_val_predict(KN1, feature1, y)
scoreKN1 = accuracy_score(y, yhat_KN1)
report = classification_report(y, yhat_KN1)
print('-----KN1 with feature1-----\n')
print(report)
print('accuracy = ' + str(scoreKN1))
```

-----KN1 with feature1-----

	precision	recall	f1-score	support
0.0	0.80	0.19	0.31	1000
1.0	0.54	0.95	0.69	1000
avg / total	0.67	0.57	0.50	2000

accuracy = 0.572

```
In [96]: RF1 = RandomForestClassifier()
yhat_RF1 = cross_val_predict(RF1, feature1, y)
scoreRF1 = accuracy_score(y, yhat_RF1)
report = classification_report(y, yhat_RF1)
print('-----RF1 with feature1-----\n')
print(report)
print('accuracy = ' + str(scoreRF1))
```

-----RF1 with feature1-----

	precision	recall	f1-score	support
0.0	0.62	0.77	0.69	1000
1.0	0.70	0.53	0.60	1000
avg / total	0.66	0.65	0.65	2000

accuracy = 0.6515

```
In [151]: SVM1 = SVC()
yhat_SVM1 = cross_val_predict(SVM1,feature1,y)
scoreSVM1 = accuracy_score(y,yhat_SVM1)
report = classification_report(y, yhat_SVM1)
print('-----SVM1 with feature1-----\n')
print(report)
print('accuracy = ' + str(scoreSVM1))
```

```
-----SVM1 with feature1-----
              precision    recall  f1-score   support

     0.0         0.80      0.70      0.74       1000
     1.0         0.73      0.82      0.77       1000

 avg / total         0.76      0.76      0.76       2000

accuracy = 0.7585
```

```
In [10]: MLP1 = MLPClassifier()
yhat_MLP1 = cross_val_predict(MLP1,feature1,y)
scoreMLP1 = accuracy_score(y,yhat_MLP1)
report = classification_report(y, yhat_MLP1)
print('-----MLP1 with feature1-----\n')
print(report)
print('accuracy = ' + str(scoreMLP1))
```

```
-----MLP1 with feature1-----
              precision    recall  f1-score   support

     0.0         0.84      0.82      0.83       1000
     1.0         0.82      0.84      0.83       1000

 avg / total         0.83      0.83      0.83       2000

accuracy = 0.8275
```

2.5 Model Tuning

Can you try to beat the simple model you created above? Here are some things you may try:

- When creating TfidfVectorizer object, you may tweak sublinear_tf parameter which use the tf with logarithmic scale instead of the usual tf.
- You may also exclude words that are too frequent or too rare, by adjusting max_df and min_df.
- Adjusting parameters available in the model, like neural network structure or number of trees in the forest.

Design at least 3 experiments using these techniques. Show your experimental results.

KNeighbors

กำหนด max_df = 0.85 min_df = 0.1

1. accuracy: 0.676 > 0.572
2. precision: 0.68 > 0.67
3. recall: 0.68 > 0.57

เปลี่ยน hyper parameter ค่า k = 300 จาก Default = 5 โดย max_df = 0.85 min_df = 0.1

1. accuracy: 0.727
2. precision: 0.74
3. recall: 0.73

```
In [58]: Tfidf2 = TfidfVectorizer(max_df = 0.85,min_df = 0.1)
feature2 = Tfidf2.fit_transform(raw_data)
yhat_KN2 = cross_val_predict(KN1,feature2,y)
scoreKN2 = accuracy_score(y,yhat_KN2)
report = classification_report(y, yhat_KN2)
print('-----KN1 with feature2-----\n')
print(report)
print('accuracy = ' + str(scoreKN2))
```

```
-----KN1 with feature2-----
              precision    recall  f1-score   support

    0.0         0.71     0.60     0.65       1000
    1.0         0.65     0.75     0.70       1000

avg / total         0.68     0.68     0.67       2000

accuracy = 0.676
```

```
In [82]: KN2 = KNeighborsClassifier(n_neighbors = 300)
yhat_KN2 = cross_val_predict(KN2,feature2,y)
scoreKN2 = accuracy_score(y,yhat_KN2)
report = classification_report(y, yhat_KN2)
print('-----KN2 with feature2-----\n')
print(report)
print('accuracy = ' + str(scoreKN2))
```

```
-----KN2 with feature2-----
              precision    recall  f1-score   support

    0.0         0.69     0.84     0.75       1000
    1.0         0.79     0.61     0.69       1000

avg / total         0.74     0.73     0.72       2000

accuracy = 0.727
```

RandomForest

กำหนด max_df = 0.85 min_df = 0.1

1. accuracy: 0.6775 > 0.6515
2. precision: 0.68 > 0.66
3. recall: 0.68 > 0.65

เปลี่ยน hyper parameter ค่า n_estimators = 300 จาก Default = 10 โดย max_df = 0.85 min_df = 0.1

1. accuracy: 0.7855
2. precision: 0.79
3. recall: 0.79

```
In [145]: Tfidf2 = TfidfVectorizer(max_df = 0.85,min_df = 0.1)
feature2 = Tfidf2.fit_transform(raw_data)
yhat_RF2 = cross_val_predict(RF1,feature2,y)
scoreRF2 = accuracy_score(y,yhat_RF2)
report = classification_report(y, yhat_RF2)
print('-----RF1 with feature2-----\n')
print(report)
print('accuracy = ' + str(scoreRF2))
```

```
-----RF1 with feature2-----
              precision    recall  f1-score   support

    0.0         0.65     0.77     0.71       1000
    1.0         0.72     0.58     0.64       1000

avg / total         0.68     0.68     0.67       2000

accuracy = 0.6775
```

```
In [149]: RF2 = RandomForestClassifier(n_estimators = 1000)
yhat_RF2 = cross_val_predict(RF2, feature2, y)
scoreRF2 = accuracy_score(y, yhat_RF2)
report = classification_report(y, yhat_RF2)
print('-----RF2 with feature2-----\n')
print(report)
print('accuracy = ' + str(scoreRF2))
```

```
-----RF1 with feature2-----
              precision    recall  f1-score   support

    0.0         0.78        0.80        0.79        1000
    1.0         0.79        0.77        0.78        1000

avg / total         0.79        0.79        0.79        2000

accuracy = 0.7855
```

SVM

กำหนด max_df = 0.85 min_df = 0.1

```
1. accuracy: 0.7285 < 0.7585
2. precision: 0.73 < 0.76
3. recall: 0.73 < 0.76
**max_df และ min_df ส่งผลต่อ performance ดังนี้
```

เปลี่ยน hyper parameter ค่า kernel = 'sigmoid' จาก Default = 'rbf' โดย max_df, min_df (default)

```
1. accuracy: 0.7585
2. precision: 0.76
3. recall: 0.76
**ผลลัพธ์ไม่ต่าง
```

เปลี่ยน hyper parameter ค่า kernel = 'linear' จาก Default = 'rbf' โดย max_df, min_df (default)

```
1. accuracy: 0.845
2. precision: 0.85
3. recall: 0.84
**ผลลัพธ์ดีขึ้น
```

```
In [159]: Tfidf2 = TfidfVectorizer(max_df = 0.85 ,min_df = 0.1 )
feature2 = Tfidf2.fit_transform(raw_data)
yhat_SVM2 = cross_val_predict(SVM1, feature2, y)
scoreSVM2 = accuracy_score(y, yhat_SVM2)
report = classification_report(y, yhat_SVM2)
print('-----SVM1 with feature2-----\n')
print(report)
print('accuracy = ' + str(scoreSVM2))
```

```
-----SVM1 with feature2-----
              precision    recall  f1-score   support

    0.0         0.77        0.66        0.71        1000
    1.0         0.70        0.80        0.75        1000

avg / total         0.73        0.73        0.73        2000

accuracy = 0.7285
```

```
In [160]: SVM2 = SVC(kernel = 'sigmoid')
yhat_SVM2 = cross_val_predict(SVM2, feature1, y)
scoreSVM2 = accuracy_score(y, yhat_SVM2)
report = classification_report(y, yhat_SVM2)
print('-----SVM2 with feature1-----\n')
print(report)
print('accuracy = ' + str(scoreSVM2))
```

```
-----SVM1 with feature2-----
              precision    recall  f1-score   support

    0.0         0.80        0.70        0.74        1000
    1.0         0.73        0.82        0.77        1000

avg / total         0.76        0.76        0.76        2000

accuracy = 0.7585
```



```
In [161]: SVM2 = SVC(kernel = 'linear')
yhat_SVM2 = cross_val_predict(SVM2, feature1, y)
scoreSVM2 = accuracy_score(y, yhat_SVM2)
report = classification_report(y, yhat_SVM2)
print('-----SVM2 with feature1-----\n')
print(report)
print('accuracy = ' + str(scoreSVM2))
```

```
-----SVM1 with feature2-----
              precision    recall  f1-score   support

    0.0         0.85        0.84        0.84         1000
    1.0         0.84        0.85        0.85         1000

avg / total         0.85        0.84        0.84         2000

accuracy = 0.845
```

MLP

กำหนด max_df = 0.85 min_df = 0.1

```
1. accuracy: 0.7655 < 0.8275
2. precision: 0.77 < 0.83
3. recall: 0.77 < 0.83
**max_df และ min_df ส่งผลต่อ performance อย่างไร
```

เปลี่ยน hyper parameter ค่า hidden_layer_sizes = (200,) จาก Default = (100,) โดย max_df, min_df (default)

```
1. accuracy: 0.828
2. precision: 0.83
3. recall: 0.83
**accuracy เพิ่มขึ้น ผลต่างขึ้นกับค่าอื่น
```

เปลี่ยน hyper parameter ค่า hidden_layer_sizes = (1000,) จาก Default = (100,) โดย max_df, min_df (default)

```
1. accuracy: 0.829
2. precision: 0.83
3. recall: 0.83
**accuracy เพิ่มขึ้น ผลต่างขึ้นกับค่าอื่น
```

```
In [32]: yhat_MLP2 = cross_val_predict(MLP1, feature2, y)
scoreMLP2 = accuracy_score(y, yhat_MLP2)
report = classification_report(y, yhat_MLP2)
print('-----MLP1 with feature2-----\n')
print(report)
print('accuracy = ' + str(scoreMLP2))
```

```
C:\Users\dell\Anaconda3\lib\site-packages\sklearn\network\multilayer_perceptron.py:564: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
% self.max_iter, ConvergenceWarning)
C:\Users\dell\Anaconda3\lib\site-packages\sklearn\network\multilayer_perceptron.py:564: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
% self.max_iter, ConvergenceWarning)
```

```
-----MLP1 with feature2-----
              precision    recall  f1-score   support

    0.0         0.77        0.76        0.76         1000
    1.0         0.76        0.77        0.77         1000

avg / total         0.77        0.77        0.77         2000

accuracy = 0.7655
```

```
C:\Users\dell\Anaconda3\lib\site-packages\sklearn\network\multilayer_perceptron.py:564: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
% self.max_iter, ConvergenceWarning)
```

```
In [168]: MLP2 = MLPClassifier(hidden_layer_sizes= (200,))
yhat_MLP2 = cross_val_predict(MLP2,feature1,y)
scoreMLP2 = accuracy_score(y,yhat_MLP2)
report = classification_report(y, yhat_MLP2)
print('-----MLP2 with feature1-----\n')
print(report)
print('accuracy = ' + str(scoreMLP2))
```

```
-----MLP1 with feature2-----
              precision    recall  f1-score   support

    0.0         0.84        0.81        0.83        1000
    1.0         0.82        0.84        0.83        1000

 avg / total         0.83        0.83        0.83        2000

accuracy = 0.828
```

```
In [169]: MLP2 = MLPClassifier(hidden_layer_sizes= (1000,))
yhat_MLP2 = cross_val_predict(MLP2,feature1,y)
scoreMLP2 = accuracy_score(y,yhat_MLP2)
report = classification_report(y, yhat_MLP2)
print('-----MLP2 with feature1-----\n')
print(report)
print('accuracy = ' + str(scoreMLP2))
```

```
-----MLP2 with feature1-----
              precision    recall  f1-score   support

    0.0         0.84        0.81        0.83        1000
    1.0         0.82        0.84        0.83        1000

 avg / total         0.83        0.83        0.83        2000

accuracy = 0.829
```

3 Text Clustering

We have heard about Google News clustering. In this exercise, we are going to implement it with Python.

3.1 Data Preprocessing

Let's switch up and use another dataset called 20newsgroup data, which is a collection of approximately 20,000 newsgroup documents, partitioned (nearly) evenly across 20 different newsgroups. The data is collected from a university's mailing list, where students exchange opinions in everything from motorcycles to middle east politics.

1. Import data using `sklearn.datasets.fetch_20newsgroups`
2. Transform data to vector with `TfidfVectorizer`

```
In [171]: from sklearn.datasets import fetch_20newsgroups
```

```
In [172]: raw_data = fetch_20newsgroups(subset = 'train') #select data in subset 'train'
```

```
In [200]: # transform data
Tfidf_google = TfidfVectorizer(max_df = 0.8 , min_df = 0.05)
x = Tfidf_google.fit_transform(raw_data)
```

3.2 Clustering

We are going to use the simplest clustering model, k-means clustering, to do this task. Our hope is that this simple algorithm will result in meaningful news categories, without using labels.

1. Fit K-Means clustering model to the text vector. What is the value of K you should pick? Why?
2. Use Silhouette score to evaluate your clusters. Try to evaluate the model for different values of k to see which k fits best for the dataset.

KMeans

select k = 20 เพราะ data ที่ import เข้ามามี 20 groups

Result

1. K = 20 Silhoutte score = 0.00307
2. K = 15 Silhoutte score = 0.00484
3. K = 10 Silhoutte score = -0.00039
4. K = 5 Silhoutte score = 0.00411
5. K = 35 Silhoutte score = 0.00525
6. K = 50 Silhoutte score = 0.00913

จากผลลัพธ์ K = 50 ดีที่สุด แต่ผมคิดว่ามันไม่ควรจะแบ่งได้ 50 cluster จึงเลือก K = 15 เพราะใน data จริงๆ จาก 20 group มันสามารถจัดกลุ่มความคล้ายกันทำให้ group น้อยลงได้

```
In [201]: from sklearn.cluster import KMeans
          from sklearn.metrics import silhouette_score
```

```
In [202]: kmean = KMeans(n_clusters = 20)
          kmean.fit(x)
          sil_score = silhouette_score(x,kmean.labels_ )
          print('-----K = 20-----')
          print('silhouette score = ' + str(sil_score))

          -----K = 20-----
          silhouette score = 0.003071184841869602
```

```
In [203]: kmean1 = KMeans(n_clusters = 15)
          kmean1.fit(x)
          sil_score = silhouette_score(x,kmean1.labels_ )
          print('-----K = 15-----')
          print('silhouette score = ' + str(sil_score))

          -----K = 15-----
          silhouette score = 0.004844898597086147
```

```
In [204]: kmean2 = KMeans(n_clusters = 10)
          kmean2.fit(x)
          sil_score = silhouette_score(x,kmean2.labels_ )
          print('-----K = 10-----')
          print('silhouette score = ' + str(sil_score))

          -----K = 10-----
          silhouette score = -0.0003900338875452058
```

```
In [205]: kmean3 = KMeans(n_clusters = 5)
          kmean3.fit(x)
          sil_score = silhouette_score(x,kmean3.labels_ )
          print('-----K = 5-----')
          print('silhouette score = ' + str(sil_score))

          -----K = 5-----
          silhouette score = 0.004115487009679443
```

```
In [206]: kmean4 = KMeans(n_clusters = 35)
          kmean4.fit(x)
          sil_score = silhouette_score(x,kmean4.labels_ )
          print('-----K = 35-----')
          print('silhouette score = ' + str(sil_score))

          -----K = 35-----
          silhouette score = 0.005255762323989831
```

```
In [207]: kmean5 = KMeans(n_clusters = 50)
          kmean5.fit(x)
          sil_score = silhouette_score(x,kmean5.labels_ )
          print('-----K = 50-----')
          print('silhouette score = ' + str(sil_score))

          -----K = 50-----
          silhouette score = 0.009132529670051202
```


3.3 Topic Terms

We want to explore each cluster to understand what news articles are in the cluster, what terms are associated with the cluster. This will require a bit of hacking.

1. Use `TfidfVectorizer.get_feature_names()` to extract words associated with each dimension of the text vector.
2. Extract cluster's centroids using `kmeans.cluster_centers_`.
3. For each centroid, print the top 15 words that have the highest frequency.

Kmeans clustering K = 15

TOP 15 word in each clusters

1. Cluster 0: her she movie you was they when their like about there up so out which
2. Cluster 1: movie was you they what just there horror me my like so about or good
3. Cluster 2: house tom horror movie you they was what out good just there her up some
4. Cluster 3: him life was story more movie which its her we their there when man into
5. Cluster 4: jack her ship movie they she him good there you james up more than love
6. Cluster 5: you movie your if out about or so her just what there they was when
7. Cluster 6: alien ship they movie was which her earth more so its than dr there when
8. Cluster 7: they movie there their up Joe you out which like comedy funny when can was
9. Cluster 8: mr robin which was her no characters there would out while more films character series
10. Cluster 9: action movie van you was some like plot so there they good out bad scenes
11. Cluster 10: war private men they movie their battle was british american world most mission action about
12. Cluster 11: killer you horror murder movie more her was its about what up they too which
13. Cluster 12: star effects planet science special fiction movie was space they series so like some we
14. Cluster 13: they you we so was about like their there movie what or just up out
15. Cluster 14: 10 you movie was me also about little some plot my well did see out

```
In [208]: feat = Tfidf_google.get_feature_names()
```

```
In [214]: cluster_centers = kmean1.cluster_centers_
```

```
In [215]: order_centroids = cluster_centers.argsort()[::-1]
for i in range(len(cluster_centers)):
    print("Cluster %d:" % i),
    for ind in order_centroids[i, :15]:
        print(' %s' % feat[ind])
```

Cluster 0:

her
she
movie
you
was
they
when
their
like
about
there
up
so
out
which

Cluster 1:

movie
was
you
they
what
just
there
horror
me
my
like
so
about
or
good

Cluster 2:

house
tom
horror
movie
you
they
was
what
out
good
just
there
her
up
some

Cluster 3:

him
life
was
story
more
movie
which
its
her
we
their
there
when
man
into

Cluster 4:

jack
her
ship
movie
they
she
him
good
there
you
james
up
more
than
love

Cluster 5:

you
movie
your
if
out
about
or
so
her
just
what
there
they
was
when

Cluster 6:

alien
ship
they
movie
was
which
her
earth
more
so
its
than
dr
there
when

Cluster 7:

they
movie
there
their
up
joe
you
out
which
like
comedy
funny
when
can
was

Cluster 8:

mr
robin
which
was
her
no
characters
there
would
out
while
more
films
character
series

Cluster 9:

action
movie
van
you
was
some
like
plot
so
there
they
good
out
bad
scenes

Cluster 10:

war
private
men
they
movie
their
battle
was
british
american
world
most
mission
action
about

Cluster 11:

killer
you
horror
murder
movie
more
her
was
its
about
what
up
they
too
which

Cluster 12:

star
effects
planet
science
special
fiction
movie
was
space
they
series
so
like
some
we

Cluster 13:

they
you
we
so
was
about
like
their
there
movie
what
or
just
up
out

Cluster 14:

10
you
movie
was
me
also
about
little
some
plot
my
well
did
see
out